INDUS11 — First Review Run Sheet

AI Fraud Decision Engine · Joshi Om · Krish Gajera · Drashti Dedaniya

Format: presentation & briefing only — no live run. Deck: **Indus11-Mentor-Deck.pptx** (6 slides).

WHO SPEAKS WHAT

#	SLIDE	SPEAKER	SAY THIS
1	Title	**Joshi (opens)**	"We're Joshi, Krish & Drashti. Indus11 decides in real time if a transaction is fraud — and explains why in plain English."
2	Idea & domain	**Joshi**	Domain = fraud detection. Problem (ms decisions; rules miss new fraud, AI can't explain) → 3 engines, one 0–100 score → APPROVE / REVIEW / BLOCK.
3	What we'll build	**Krish**	Four deliverables: analysis API, live dashboard, synthetic dataset, explained decisions. Scope = working prototype on synthetic data.
4	Tools & technology	**All — each own tools**	Joshi: FastAPI · MongoDB · Redis · Docker. Krish: Neo4j. Drashti: ChromaDB · LangChain + Ollama · React.
5	AI integration	**Drashti**	RAG flow: flagged → retrieve patterns (ChromaDB) → LLM (Ollama llama3) reasons → score + explanation. Capped 30/100, temperature 0.
6	Team & contributions	**All — each reads own column**	Each member states their layer & contributions — this is the task distribution.

Q&A QUICK HITS

Is it actually working? Yes — 8/8 tests pass, it runs end-to-end, and we've analyzed real transactions. Happy to demo live next time.

Why MongoDB? Flexible documents, no migrations; store/query records by field.

Why Neo4j? Fraud rings are relationships — one-line graph query vs painful recursive SQL.

Why 4 databases / not one? Right tool per job — graph, cache, vectors are different problems. Single-DB option would be Postgres + pgvector, but graph traversal suffers.

How is AI integrated? RAG — retrieve similar patterns from ChromaDB; LLM (Ollama llama3) scores & explains; capped at 30/100.

How do you know it works? 8 tests + we plant fraud rings in synthetic data and confirm detection. Next: precision/recall.

OPENERS & CLOSERS

Open: "Indus11 is a real-time fraud engine that scores every transaction 0–100 and returns APPROVE, REVIEW or BLOCK — with a plain-English reason."

Hand-off: each member names their layer before speaking to it ("I own the graph side, so…").

On scope: "It's a working prototype on synthetic data — proves the approach end-to-end, not a production bank."

Close: "Most of the build is done; next is measuring accuracy and hardening. Questions welcome."

REMEMBER

Thresholds: **APPROVE 0–39 · REVIEW 40–69 · BLOCK 70–100**.

Don't say "the AI does it" — say *how* (retrieve → reason → explain).

If you don't know something, say so + how you'd find out.

Goal of this review: convey the idea, the technology choices, the task split, and the plan — clearly and accurately. No system needs to run.